

6.824 2024 Midterm Exam: Crash Recovery in File Systems

Time: 120 minutes Instructions: - You may consult 6.824 lecture notes, papers, and lab code. - If you have clarifying questions, ask them in a **private** Piazza post to the staff. - Do **not** discuss the exam with anyone. - **Show all your work** for full credit. Partial credit will be given for reasoning, even if the final answer is incorrect. - **Assume** a UNIX-like file system unless otherwise specified.

Section A: Short Answer (30 points)

Answer each question concisely (1–3 sentences).

1. (5 pts) The slides describe a "delete/create inconsistency" where an `unlink("f1")` followed by a `create("f2")` can lead to undetected corruption after a crash. Briefly explain **why** this inconsistency is *undetected* by `fsck`.
 2. (5 pts) What is the **key invariant** that UNIX file systems aim to preserve during crash recovery regarding directory entries and i-nodes? State it precisely.
 3. (5 pts) Why does the UNIX `fsck` program sometimes ask the user for input during recovery? Give an example scenario where this might happen.
 4. (5 pts) The slides mention that synchronous writes in file creation can be reduced by deferring some operations. Which **two writes** (from the 7-step `create` sequence) are **most critical** to perform synchronously, and why?
 5. (5 pts) Explain why a write-back disk cache improves performance for file creation but introduces crash recovery challenges. Use the `create("d/f")` example from the slides in your answer.
 6. (5 pts) Suppose a file system **never** writes the free block bitmap to disk (step #1 and #5 in `create`). How would `fsck` recover the correct free list after a crash? Would this approach work for **all** crash scenarios? Justify your answer.
-

Section B: Problem Solving (40 points)

Work through each scenario carefully. Show your reasoning.

B1. Crash Recovery Scenarios (20 pts)

Consider the following sequence of operations on a UNIX file system with a write-back cache: `c`
`fd = create("dir/file1", 0666); write(fd, "abcde", 5);`
`unlink("dir/file1"); fd = create("dir/file2", 0666); write(fd,`
`"fghij", 5);` The disk writes for these operations (from the slides) are: -
`create("file1")`: 1, 2, 3, 4, 5, 6, 7 - `unlink("file1")`: 8, 9, 10, 11 -
`create("file2")`: 1, 2, 3, 4, 5, 6, 7

Assume the cache flushes **only the following writes** to disk before a crash: - For `file1`: writes 1, 2, 3, 6 - For `unlink("file1")`: write 8 - For `file2`: writes 1, 2, 3

After recovery: 1. (5 pts) What does the file system tree look like? List the visible files in `dir/` and their contents (if any). 2. (5 pts) Which i-nodes and blocks (if any) are **leaked** (allocated but unreachable)? How would `fsck` handle them? 3. (5 pts) Is this a **benign** outcome (i.e., could it have resulted from a crash at a slightly earlier or later time)? Why or why not? 4. (5 pts) Propose a **minimal set of synchronous writes** (from the original 7-step sequence) that would prevent this inconsistency. Justify your choice.

B2. `rename()` Crash Recovery (20 pts)

The slides briefly mention `rename()` but do not detail its crash recovery challenges. Suppose `rename("a/b", "a/c")` performs the following disk writes: 1. Update the i-node for `b` (if its name is stored there). 2. Remove the directory entry "`b`" from `a`'s contents. 3. Add the directory entry "`c`" to `a`'s contents. 4. Update `a`'s i-node (mtime, length).

Assume a crash occurs after **only step 2** completes.

1. (5 pts) What does the file system look like after recovery? Is `b` or `c` visible? What happens to `b`'s i-node and data?
 2. (5 pts) Could `fsck` automatically repair this? If not, what would it ask the user?
 3. (5 pts) Propose a **synchronous write order** for `rename()` that guarantees recovery to a consistent state. Hint: Think about commit points.
 4. (5 pts) How would your answer change if `b` and `c` were in **different directories** (e.g., `rename("a/b", "x/c")`)? List the additional challenges.
-

Section C: Design (30 points)

Apply concepts from the slides to design solutions.

C1. Optimizing `create()` (15 pts)

The slides suggest that only **two synchronous writes** (steps 2 and 6) are strictly necessary for correct `create()` recovery, but UNIX synchronously writes steps 2 and 3.

1. (5 pts) Why might UNIX choose to synchronously write step 3 (`d`'s contents) even if it's not strictly required for correctness? Hint: Consider `fsck`'s behavior.
 2. (5 pts) Suppose we **defer step 3** (as the slides suggest) but still synchronously write step 2. Describe a crash scenario where this leads to a **user-visible inconsistency** (e.g., a file appears to exist but cannot be opened).
 3. (5 pts) Propose a **hybrid approach** that reduces synchronous writes while avoiding the inconsistency in (2). Your solution may involve additional metadata or reordering.
-

C2. Journaling vs. fsck (15 pts)

The slides focus on fsck-based recovery, but modern file systems often use **journaling** (e.g., ext4, XFS).

1. (5 pts) How would a **journaling file system** (e.g., with metadata journaling) handle the `unlink("f1") + create("f2")` inconsistency from the slides? Would it detect the issue? Why or why not?
 2. (5 pts) Journaling introduces **write overhead**. For the 7-step `create()` sequence, which steps would be **journalled**, and which could be written directly to their final locations? Justify your answer.
 3. (5 pts) A critic argues: "*Journaling is unnecessary because fsck can always recover the file system to a consistent state.*" Write a **counterargument** using examples from the slides (e.g., undetected inconsistencies or performance trade-offs).
-

Section D: Essay (20 points)

Choose **one** of the following. Write a clear, structured response (1–2 paragraphs).

1. **Performance vs. Durability Trade-offs** The slides emphasize that crash recovery must be **fast** because persistent storage is slow. Discuss how the following techniques balance this trade-off:
 - Write-back caching
 - Synchronous writes for commit points
 - fsck vs. journaling Use examples from the slides (e.g., `create()` throughput) to support your argument.
 2. **Distributed File Systems** The slides note that crash recovery in disk file systems shares similarities with distributed systems. Compare and contrast:
 - How a **single-machine file system** (e.g., UNIX) handles recovery after a crash.
 - How a **distributed storage system** (e.g., Spanner or FaRM from previous exams) handles recovery after a node failure. Focus on **ordering guarantees, commit protocols, and user-visible inconsistencies**.
-

End of Exam Good luck! Remember to show all your work.